

# Food Classification with Deep Learning in Keras / Tensorflow

*Computer, what am I eating anyway?*



```
1. from IPython.display import HTML, Image
2.
3. url = 'http://stratospark.com/demos/food-101/'
4. el = '<' + 'iframe src="{0}"'.format(url) + ' width="100%" height=600></iframe>'
5. HTML(el)
```

If you are reading this on GitHub, the demo looks like this. Please follow the link below to view the live demo on my blog.

```
1. Image('demo.jpg')
```

jpeg

**Demo available @** <http://blog.stratospark.com/deep-learning-applied-food-classification-deep-learning-keras.html>

**Code available @** <https://github.com/stratospark/food-101-keras>

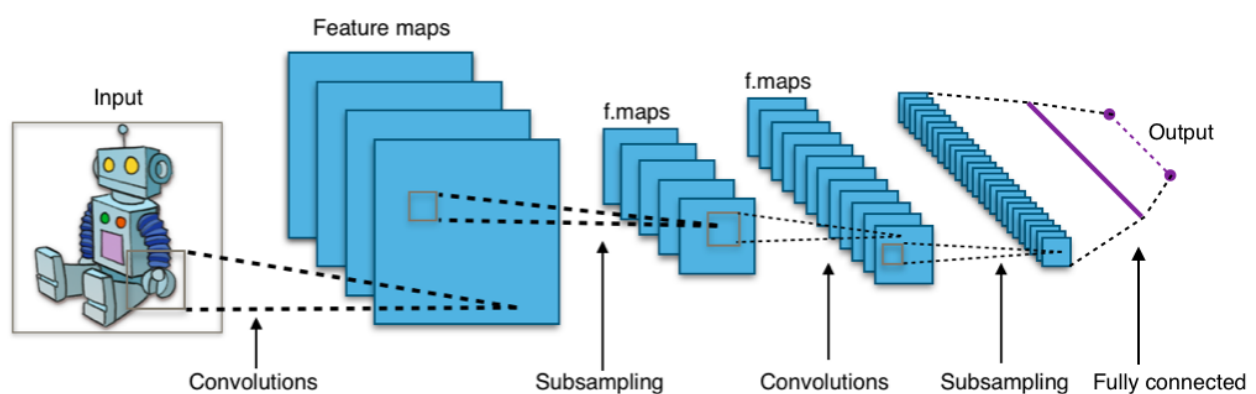
## UPDATES

- **2017-03-22** Learn how to use this model in a mobile app:  
<http://blog.stratospark.com/creating-a-deep-learning-ios-app-with-keras-and-tensorflow.html>

- Introduction
  - Project Description
  - Approach
  - Results
  - Thoughts
- Experiment
  - Loading and Preprocessing Dataset
  - Visualization Tools
  - Image Augmentation
  - Training
  - Model Evaluation
  - Results Visualization
  - Interactive Classification
  - Keras.js Export

## Introduction

Convolutional Neural Networks (CNN), a technique within the broader Deep Learning field, have been a revolutionary force in Computer Vision applications, especially in the past half-decade or so. One main use-case is that of image classification, e.g. determining whether a picture is that of a dog or cat.



You don't have to limit yourself to a binary classifier of course; CNNs can easily scale to thousands of different classes, as seen in the well-known ImageNet dataset of 1000 classes, used to benchmark computer vision algorithm performance.



In the past couple of years, these cutting edge techniques have started to become available to the broader software development community. Industrial strength packages such as [Tensorflow](#) have given us the same building blocks that Google uses to write deep learning applications for embedded/mobile devices to scalable clusters in the cloud – *Without having to handcode the GPU matrix operations, partial derivative gradients, and stochastic optimizers that make efficient applications possible.*

On top of all of this, are user-friendly APIs such as [Keras](#) that abstract away some of the lower level details and allow us to focus on rapidly prototyping a deep learning computation graph. Much like we would mix and match Legos to get a desired result.

## Project Description

As an introductory project for myself, I chose to use a pre-trained image classifier that comes with Keras, and retrain it on a dataset that I find interesting. I'm very much into

good food and home cooking, so something along those lines was appetizing.

In the paper, [Food-101 – Mining Discriminative Components with Random Forests](#), they introduce the Food-101 dataset. There are 101 different classes of food, with 1000 labeled images per class available for supervised training.



## Approach

I was inspired by this Keras blog post: [Building powerful image classification models using very little data](#), and a related script I found on github: [keras-finetuning](#).

I built a system recently for the purpose of experimenting with Deep Learning. The key components are an Nvidia Titan X Pascal w/12 GB of memory, 96 GB of system RAM, as well as a 12-core Intel Core i7. It is running 64-bit Ubuntu 16.04 and using the Anaconda Python distribution. Unfortunately, you won't be able to follow along with this notebook on your own system unless you have enough RAM. In the future, I would like to learn how to handle larger than RAM datasets in a performant way. **Please get in touch if you have any ideas!**

I've spent about 1 month on and off building this project, trying to train dozens of models and exploring various areas such as multiprocessing for faster image augmentation. This is a cleaned up version of the notebook that contains my best performing model as of Jan 22, 2017.

## Results

After fine-tuning a pre-trained Google [InceptionV3](#) model, I was able to achieve about **82.03% Top-1 Accuracy** on the test set using a single crop per item. Using 10 crops per example and taking the most frequent predicted class(es), I was able to achieve **86.97% Top-1 Accuracy** and **97.42% Top-5 Accuracy**

Others have been able to achieve more accurate results:

- **InceptionV3: 88.28% Top-1 Accuracy** with unknown-crops. [Hassannejad, Hamid, et al. "Food Image Recognition Using Very Deep Convolutional Networks." Proceedings of the 2nd International Workshop on Multimedia Assisted Dietary Management. ACM, 2016.](#)

- **ResNet200: 90.14% Top-1 Accuracy** on the Food-101 dataset augmented with 19 Korean dishes. [NVIDIA DEEP LEARNING CONTEST 2016](#), Keun-dong Lee, DaUn Jeong, Seungjae Lee, Hyung Kwan Son (ETRI VisualBrowsing Team), Oct.7, 2016.
- **WiSeR: 90.27% Top-1 Accuracy** with 10-crops. [Martinel, Niki, Gian Luca Foresti, and Christian Micheloni. "Wide-Slice Residual Networks for Food Recognition." arXiv preprint arXiv:1612.06543 \(2016\).](#)

## Thoughts

- Loading a large amount of data into memory, how to avoid?
  - Saving the data into h5py file for out of band processing?
  - Using Dask for distributed processing?
  - Improving multiprocessing image augmentation?
- 

- Exporting to Tensorflow [mobile app](#)?

Implemented! Check out: <http://blog.stratospark.com/creating-a-deep-learning-ios-app-with-keras-and-tensorflow.html>

## Experiment

### Loading and Preprocessing Dataset

Let's import all of the packages needed for the rest of the notebook:

```
1. import matplotlib.pyplot as plt
2. import matplotlib.image as img
3. import numpy as np
4. from scipy.misc import imresize
5.
6. %matplotlib inline
7.
8. import os
9. from os import listdir
10. from os.path import isfile, join
11. import shutil
12. import stat
13. import collections
14. from collections import defaultdict
15.
16. from ipywidgets import interact, interactive, fixed
17. import ipywidgets as widgets
18.
19. import h5py
20. from sklearn.model_selection import train_test_split
21. from keras.utils.np_utils import to_categorical
22. from keras.applications.inception_v3 import preprocess_input
23. from keras.models import load_model
```

Using TensorFlow backend.

Download the dataset and extract it within the notebook folder. It may be easier to do this in a separate terminal window.

```
1. # !wget http://data.vision.ee.ethz.ch/cvl/food-101.tar.gz
```

```
1. # !tar xzvf food-101.tar.gz
```

Let's see what sort of foods are represented here:

```
1. !ls food-101/images
```

apple_pie	eggs_benedict	onion_rings
baby_back_ribs	escargots	oysters
baklava	falafel	pad_thai
beef_carpaccio	filet_mignon	paella
beef_tartare	fish_and_chips	pancakes
beet_salad	foie_gras	panna_cotta
beignets	french_fries	peking_duck
bibimbap	french_onion_soup	pho
bread_pudding	french_toast	pizza
breakfast_burrito	fried_calamari	pork_chop
bruschetta	fried_rice	poutine
caesar_salad	frozen_yogurt	prime_rib
cannoli	garlic_bread	pulled_pork_sandwich
caprese_salad	gnocchi	ramen
carrot_cake	greek_salad	ravioli
ceviche	grilled_cheese_sandwich	red_velvet_cake
cheesecake	grilled_salmon	risotto
cheese_plate	guacamole	samosa
chicken_curry	gyoza	sashimi
chicken_quesadilla	hamburger	scallops
chicken_wings	hot_and_sour_soup	seaweed_salad
chocolate_cake	hot_dog	shrimp_and_grits
chocolate_mousse	huevos_rancheros	spaghetti_bolognese
churros	hummus	spaghetti_carbonara
clam_chowder	ice_cream	spring_rolls
club_sandwich	lasagna	steak
crab_cakes	lobster_bisque	strawberry_shortcake
creme_brulee	lobster_roll_sandwich	sushi
croque_madame	macaroni_and_cheese	tacos
cup_cakes	macarons	takoyaki
deviled_eggs	miso_soup	tiramisu
donuts	mussels	tuna_tartare
dumplings	nachos	waffles
edamame	omelette	

```
1. !ls food-101/images/apple_pie/ | head -10
```

```
1005649.jpg
1011328.jpg
101251.jpg
1014775.jpg
1026328.jpg
1028787.jpg
1034399.jpg
103801.jpg
1038694.jpg
1043283.jpg
ls: write error: Broken pipe
```

Let's look at some random images from each food class. You can right click and open the



image in a new window or save it in order to see it at a higher resolution.

```
1. root_dir = 'food-101/images/'
2. rows = 17
3. cols = 6
4. fig, ax = plt.subplots(rows, cols, frameon=False, figsize=(15, 25))
5. fig.suptitle('Random Image from Each Food Class', fontsize=20)
6. sorted_food_dirs = sorted(os.listdir(root_dir))
7. for i in range(rows):
8.     for j in range(cols):
9.         try:
10.            food_dir = sorted_food_dirs[i*cols + j]
11.        except:
12.            break
13.        all_files = os.listdir(os.path.join(root_dir, food_dir))
14.        rand_img = np.random.choice(all_files)
15.        img = plt.imread(os.path.join(root_dir, food_dir, rand_img))
16.        ax[i][j].imshow(img)
17.        ec = (0, .6, .1)
18.        fc = (0, .7, .2)
19.        ax[i][j].text(0, -20, food_dir, size=10, rotation=0,
20.                      ha="left", va="top",
21.                      bbox=dict(boxstyle="round", ec=ec, fc=fc))
22. plt.setp(ax, xticks=[], yticks=[])
23. plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

png

A `multiprocessing.Pool` will be used to accelerate image augmentation during training.

```
1. # Setup multiprocessing pool
2. # Do this early, as once images are loaded into memory there will be Errno 12
3. # http://stackoverflow.com/questions/14749897/python-multiprocessing-memory-u
4. import multiprocessing as mp
5.
6. num_processes = 6
7. pool = mp.Pool(processes=num_processes)
```

We need maps from class to index and vice versa, for proper label encoding and pretty printing.

```
1. class_to_ix = {}
2. ix_to_class = {}
3. with open('food-101/meta/classes.txt', 'r') as txt:
4.     classes = [l.strip() for l in txt.readlines()]
5.     class_to_ix = dict(zip(classes, range(len(classes))))
6.     ix_to_class = dict(zip(range(len(classes)), classes))
7.     class_to_ix = {v: k for k, v in ix_to_class.items()}
8. sorted_class_to_ix = collections.OrderedDict(sorted(class_to_ix.items()))
```

The Food-101 dataset has a provided train/test split. We want to use this in order to compare our classification performance with other implementations.

```

1. # Only split files if haven't already
2. if not os.path.isdir('./food-101/test') and not os.path.isdir('./food-101/train'):
3.
4.     def copytree(src, dst, symlinks = False, ignore = None):
5.         if not os.path.exists(dst):
6.             os.makedirs(dst)
7.             shutil.copystat(src, dst)
8.         lst = os.listdir(src)
9.         if ignore:
10.            excl = ignore(src, lst)
11.            lst = [x for x in lst if x not in excl]
12.        for item in lst:
13.            s = os.path.join(src, item)
14.            d = os.path.join(dst, item)
15.            if symlinks and os.path.islink(s):
16.                if os.path.lexists(d):
17.                    os.remove(d)
18.                os.symlink(os.readlink(s), d)
19.            try:
20.                st = os.lstat(s)
21.                mode = stat.S_IMODE(st.st_mode)
22.                os.lchmod(d, mode)
23.            except:
24.                pass # lchmod not available
25.            elif os.path.isdir(s):
26.                copytree(s, d, symlinks, ignore)
27.            else:
28.                shutil.copy2(s, d)
29.
30.        def generate_dir_file_map(path):
31.            dir_files = defaultdict(list)
32.            with open(path, 'r') as txt:
33.                files = [l.strip() for l in txt.readlines()]
34.                for f in files:
35.                    dir_name, id = f.split('/')
36.                    dir_files[dir_name].append(id + '.jpg')
37.            return dir_files
38.
39.        train_dir_files = generate_dir_file_map('food-101/meta/train.txt')
40.        test_dir_files = generate_dir_file_map('food-101/meta/test.txt')
41.
42.
43.        def ignore_train(d, filenames):
44.            print(d)
45.            subdir = d.split('/')[-1]
46.            to_ignore = train_dir_files[subdir]
47.            return to_ignore
48.
49.        def ignore_test(d, filenames):
50.            print(d)
51.            subdir = d.split('/')[-1]
52.            to_ignore = test_dir_files[subdir]
53.            return to_ignore

```

```
54.  
55.     copytree('food-101/images', 'food-101/test', ignore=ignore_train)  
56.     copytree('food-101/images', 'food-101/train', ignore=ignore_test)  
57.  
58. else:  
59.     print('Train/Test files already copied into separate folders.')
```

```
Train/Test files already copied into separate folders.
```

We are now ready to load the training and testing images into memory. After everything is loaded, about 80 GB of memory will be allocated.

Any images that have a width or length smaller than `min_size` will be resized. This is so that we can take proper-sized crops during image augmentation.

```

1. %%time
2.
3. # Load dataset images and resize to meet minimum width and height pixel size
4. def load_images(root, min_side=299):
5.     all_imgs = []
6.     all_classes = []
7.     resize_count = 0
8.     invalid_count = 0
9.     for i, subdir in enumerate(listdir(root)):
10.         imgs = listdir(join(root, subdir))
11.         class_ix = class_to_ix[subdir]
12.         print(i, class_ix, subdir)
13.         for img_name in imgs:
14.             img_arr = img.imread(join(root, subdir, img_name))
15.             img_arr_rs = img_arr
16.             try:
17.                 w, h, _ = img_arr.shape
18.                 if w < min_side:
19.                     wpercent = (min_side/float(w))
20.                     hsize = int((float(h)*float(wpercent)))
21.                     #print('new dims:', min_side, hsize)
22.                     img_arr_rs = imresize(img_arr, (min_side, hsize))
23.                     resize_count += 1
24.                 elif h < min_side:
25.                     hpercent = (min_side/float(h))
26.                     wsize = int((float(w)*float(hpercent)))
27.                     #print('new dims:', wsize, min_side)
28.                     img_arr_rs = imresize(img_arr, (wsize, min_side))
29.                     resize_count += 1
30.                 all_imgs.append(img_arr_rs)
31.                 all_classes.append(class_ix)
32.             except:
33.                 print('Skipping bad image: ', subdir, img_name)
34.                 invalid_count += 1
35.         print(len(all_imgs), 'images loaded')
36.         print(resize_count, 'images resized')
37.         print(invalid_count, 'images skipped')
38.     return np.array(all_imgs), np.array(all_classes)
39.
40. X_test, y_test = load_images('food-101/test', min_side=299)

```

0 41 french\_onion\_soup  
1 99 tuna\_tartare  
2 2 baklava  
3 12 cannoli  
4 8 bread\_pudding  
5 58 ice\_cream  
6 63 macarons  
7 38 fish\_and\_chips  
8 3 beef\_carpaccio  
9 59 lasagna  
10 84 risotto  
11 53 hamburger  
12 7 bibimbap  
13 15 ceviche  
14 92 spring\_rolls  
15 78 poutine  
16 76 pizza  
17 19 chicken\_quesadilla  
18 71 paella  
19 11 caesar\_salad  
20 30 deviled\_eggs  
21 40 french\_fries  
22 25 club\_sandwich  
23 77 pork\_chop  
24 31 donuts  
25 93 steak  
26 43 fried\_calamari  
27 52 gyoza  
28 20 chicken\_wings  
29 47 gnocchi  
30 46 garlic\_bread  
31 81 ramen  
32 86 sashimi  
33 100 waffles  
34 60 lobster\_bisque  
35 23 churros  
36 1 baby\_back\_ribs  
37 0 apple\_pie  
38 27 creme\_brulee  
39 79 prime\_rib  
40 54 hot\_and\_sour\_soup  
41 55 hot\_dog  
42 82 ravioli  
43 66 nachos  
44 85 samosa  
45 95 sushi  
46 70 pad\_thai  
47 87 scallops  
48 42 french\_toast  
49 13 caprese\_salad  
50 21 chocolate\_cake

51 83 red\_velvet\_cake  
52 88 seaweed\_salad  
53 96 tacos  
54 16 cheesecake  
55 90 spaghetti\_bolognese  
56 94 strawberry\_shortcake  
57 64 miso\_soup  
58 98 tiramisu  
59 74 peking\_duck  
60 17 cheese\_plate  
61 69 oysters  
62 14 carrot\_cake  
63 6 beignets  
64 61 lobster\_roll\_sandwich  
65 45 frozen\_yogurt  
66 24 clam\_chowder  
67 9 breakfast\_burrito  
68 72 pancakes  
69 32 dumplings  
70 57 hummus  
71 10 bruschetta  
72 44 fried\_rice  
73 97 takoyaki  
74 50 grilled\_salmon  
75 4 beef\_tartare  
76 89 shrimp\_and\_grits  
77 28 croque\_madame  
78 49 grilled\_cheese\_sandwich  
79 80 pulled\_pork\_sandwich  
80 56 huevos\_rancheros  
81 35 escargots  
82 91 spaghetti\_carbonara  
83 34 eggs\_benedict  
84 33 edamame  
85 22 chocolate\_mousse  
86 18 chicken\_curry  
87 65 mussels  
88 36 falafel  
89 37 filet\_mignon  
90 26 crab\_cakes  
91 48 greek\_salad  
92 5 beet\_salad  
93 51 guacamole  
94 29 cup\_cakes  
95 68 onion\_rings  
96 39 foie\_gras  
97 67 omelette  
98 73 panna\_cotta  
99 75 pho  
100 62 macaroni\_and\_cheese  
25250 images loaded

```
693 images resized
0 images skipped
CPU times: user 1min 18s, sys: 4.82 s, total: 1min 23s
Wall time: 1min 23s
```

```
1. %%time
2. X_train, y_train = load_images('food-101/train', min_side=299)
```



0 41 french\_onion\_soup  
1 99 tuna\_tartare  
2 2 baklava  
3 12 cannoli  
4 8 bread\_pudding  
Skipping bad image: bread\_pudding 1375816.jpg  
5 58 ice\_cream  
6 63 macarons  
7 38 fish\_and\_chips  
8 3 beef\_carpaccio  
9 59 lasagna  
Skipping bad image: lasagna 3787908.jpg  
10 84 risotto  
11 53 hamburger  
12 7 bibimbap  
13 15 ceviche  
14 92 spring\_rolls  
15 78 poutine  
16 76 pizza  
17 19 chicken\_quesadilla  
18 71 paella  
19 11 caesar\_salad  
20 30 deviled\_eggs  
21 40 french\_fries  
22 25 club\_sandwich  
23 77 pork\_chop  
24 31 donuts  
25 93 steak  
Skipping bad image: steak 1340977.jpg  
26 43 fried\_calamari  
27 52 gyoza  
28 20 chicken\_wings  
29 47 gnocchi  
30 46 garlic\_bread  
31 81 ramen  
32 86 sashimi  
33 100 waffles  
34 60 lobster\_bisque  
35 23 churros  
36 1 baby\_back\_ribs  
37 0 apple\_pie  
38 27 creme\_brulee  
39 79 prime\_rib  
40 54 hot\_and\_sour\_soup  
41 55 hot\_dog  
42 82 ravioli  
43 66 nachos  
44 85 samosa  
45 95 sushi  
46 70 pad\_thai  
47 87 scallops

48 42 french\_toast  
49 13 caprese\_salad  
50 21 chocolate\_cake  
51 83 red\_velvet\_cake  
52 88 seaweed\_salad  
53 96 tacos  
54 16 cheesecake  
55 90 spaghetti\_bolognese  
56 94 strawberry\_shortcake  
57 64 miso\_soup  
58 98 tiramisu  
59 74 peking\_duck  
60 17 cheese\_plate  
61 69 oysters  
62 14 carrot\_cake  
63 6 beignets  
64 61 lobster\_roll\_sandwich  
65 45 frozen\_yogurt  
66 24 clam\_chowder  
67 9 breakfast\_burrito  
68 72 pancakes  
69 32 dumplings  
70 57 hummus  
71 10 bruschetta  
72 44 fried\_rice  
73 97 takoyaki  
74 50 grilled\_salmon  
75 4 beef\_tartare  
76 89 shrimp\_and\_grits  
77 28 croque\_madame  
78 49 grilled\_cheese\_sandwich  
79 80 pulled\_pork\_sandwich  
80 56 huevos\_rancheros  
81 35 escargots  
82 91 spaghetti\_carbonara  
83 34 eggs\_benedict  
84 33 edamame  
85 22 chocolate\_mousse  
86 18 chicken\_curry  
87 65 mussels  
88 36 falafel  
89 37 filet\_mignon  
90 26 crab\_cakes  
91 48 greek\_salad  
92 5 beet\_salad  
93 51 guacamole  
94 29 cup\_cakes  
95 68 onion\_rings  
96 39 foie\_gras  
97 67 omelette  
98 73 panna\_cotta

```
99 75 pho
100 62 macaroni_and_cheese
75747 images loaded
2091 images resized
3 images skipped
CPU times: user 3min 51s, sys: 13.9 s, total: 4min 5s
Wall time: 4min 5s
```

```
1. print('X_train shape', X_train.shape)
2. print('y_train shape', y_train.shape)
3. print('X_test shape', X_test.shape)
4. print('y_test shape', y_test.shape)
```

```
X_train shape (75747,)
y_train shape (75747,)
X_test shape (25250,)
y_test shape (25250,)
```

## Visualization Tools

```
1. @interact(n=(0, len(X_train)))
2. def show_pic(n):
3.     plt.imshow(X_train[n])
4.     print('class:', y_train[n], ix_to_class[y_train[n]])
```

```
class: 21 chocolate_cake
```

png

```
1. @interact(n=(0, len(X_test)))
2. def show_pic(n):
3.     plt.imshow(X_test[n])
4.     print('class:', y_test[n], ix_to_class[y_test[n]])
```

```
class: 21 chocolate_cake
```

png

```

1. @interact(n_class=sorted_class_to_ix)
2. def show_random_images_of_class(n_class=0):
3.     print(n_class)
4.     nrows = 4
5.     ncols = 8
6.     fig, axes = plt.subplots(nrows=nrows, ncols=ncols)
7.     fig.set_size_inches(12, 8)
8.     #fig.tight_layout()
9.     imgs = np.random.choice((y_train == n_class).nonzero()[0], nrows * ncols)
10.    for i, ax in enumerate(axes.flat):
11.        im = ax.imshow(X_train[imgs[i]])
12.        ax.set_axis_off()
13.        ax.title.set_visible(False)
14.        ax.xaxis.set_ticks([])
15.        ax.yaxis.set_ticks([])
16.        for spine in ax.spines.values():
17.            spine.set_visible(False)
18.    plt.subplots_adjust(left=0, wspace=0, hspace=0)
19.    plt.show()

```

0

png

```

1. @interact(n_class=sorted_class_to_ix)
2. def show_random_images_of_class(n_class=0):
3.     print(n_class)
4.     nrows = 4
5.     ncols = 8
6.     fig, axes = plt.subplots(nrows=nrows, ncols=ncols)
7.     fig.set_size_inches(12, 8)
8.     #fig.tight_layout()
9.     imgs = np.random.choice((y_test == n_class).nonzero()[0], nrows * ncols)
10.    for i, ax in enumerate(axes.flat):
11.        im = ax.imshow(X_test[imgs[i]])
12.        ax.set_axis_off()
13.        ax.title.set_visible(False)
14.        ax.xaxis.set_ticks([])
15.        ax.yaxis.set_ticks([])
16.        for spine in ax.spines.values():
17.            spine.set_visible(False)
18.    plt.subplots_adjust(left=0, wspace=0, hspace=0)
19.    plt.show()

```

0

png

## Image Augmentation

We need to one-hot encode each label value to create a vector of binary features rather than one feature that can take on `n_classes` values.

```
1. from keras.utils.np_utils import to_categorical
2.
3. n_classes = 101
4. y_train_cat = to_categorical(y_train, nb_classes=n_classes)
5. y_test_cat = to_categorical(y_test, nb_classes=n_classes)
```

```
1. from keras.applications.inception_v3 import InceptionV3
2. from keras.applications.inception_v3 import preprocess_input, decode_predictions
3. from keras.preprocessing import image
4. from keras.layers import Input
5.
6. import tools.image_gen_extended as T
7.
8. # Useful for checking the output of the generators after code change
9. #from importlib import reload
10. #reload(T)
```

I needed to have a more powerful Image Augmentation pipeline than the one that ships with Keras. Luckily, I was able to find this [modified version](#) to use as my base.

The author had added an extensible pipeline, which made it possible to specify additional modifications such as custom cropping functions and being able to use the Inception image preprocessor. Being able to apply preprocessing dynamically was necessary, as I did not have enough memory to keep all of the training set as `float32s`. I was able to load the entire training set as `uint8s`.

Furthermore, I was not fully utilizing either my GPU or my multicore CPU. By default, Python is only able to use a single core, thereby limiting the amount of processed/augmented images I could send to the GPU for training. Based on some performance monitoring, I was only using a small percentage of the GPU on average. By incorporating a python `multiprocessing Pool`, I was able to get about 50% CPU utilization and 90% GPU utilization.

The end result is that each epoch of training went from 45 minutes to 22 minutes! You can run the GPU graphs yourselves while training in this [notebook](#). The inspiration for trying to improve data augmentation and GPU performance came from [Jimmie Goode: Buffered Python generators for data augmentation](#)

At the moment, the code is fairly buggy and requires restarting the Python kernel whenever training is manually interrupted. The code is quite hacked together and certain features, like those that involve fitting, are disabled. I hope to improve this ImageDataGenerator and release it to the community in the future.

```
1. display(Image('./gpu.png'))
```

png

```
1. %%time
2.
3. # this is the augmentation configuration we will use for training
4. train_datagen = T.ImageDataGenerator(
5.     featurewise_center=False, # set input mean to 0 over the dataset
6.     samplewise_center=False, # set each sample mean to 0
7.     featurewise_std_normalization=False, # divide inputs by std of the dataset
8.     samplewise_std_normalization=False, # divide each input by its std
9.     zca_whitening=False, # apply ZCA whitening
10.    rotation_range=0, # randomly rotate images in the range (degrees, 0 to 180)
11.    width_shift_range=0.2, # randomly shift images horizontally (fraction of total width)
12.    height_shift_range=0.2, # randomly shift images vertically (fraction of total height)
13.    horizontal_flip=True, # randomly flip images
14.    vertical_flip=False, # randomly flip images
15.    zoom_range=[.8, 1],
16.    channel_shift_range=30,
17.    fill_mode='reflect')
18. train_datagen.config['random_crop_size'] = (299, 299)
19. train_datagen.set_pipeline([T.random_transform, T.random_crop, T.preprocess_input])
20. train_generator = train_datagen.flow(X_train, y_train_cat, batch_size=64, seed=11)
```

```
1. test_datagen = T.ImageDataGenerator()
2. test_datagen.config['random_crop_size'] = (299, 299)
3. test_datagen.set_pipeline([T.random_transform, T.random_crop, T.preprocess_input])
4. test_generator = test_datagen.flow(X_test, y_test_cat, batch_size=64, seed=11)
```

We can see what sorts of images are coming out of these ImageDataGenerators:

```
1. def reverse_preprocess_input(x0):
2.     x = x0 / 2.0
3.     x += 0.5
4.     x *= 255.
5.     return x
```

```

1. %%time
2. @interact()
3. def show_images(unprocess=True):
4.     for x in test_generator:
5.         fig, axes = plt.subplots(nrows=8, ncols=4)
6.         fig.set_size_inches(8, 8)
7.         page = 0
8.         page_size = 32
9.         start_i = page * page_size
10.        for i, ax in enumerate(axes.flat):
11.            img = x[0][i+start_i]
12.            if unprocess:
13.                im = ax.imshow( reverse_preprocess_input(img).astype('uint8'))
14.            else:
15.                im = ax.imshow(img)
16.            ax.set_axis_off()
17.            ax.title.set_visible(False)
18.            ax.xaxis.set_ticks([])
19.            ax.yaxis.set_ticks([])
20.            for spine in ax.spines.values():
21.                spine.set_visible(False)
22.
23.        plt.subplots_adjust(left=0, wspace=0, hspace=0)
24.        plt.show()
25.        break

```

png

```

CPU times: user 1.54 s, sys: 524 ms, total: 2.06 s
Wall time: 2.24 s

```

```

1. %%time
2. show_images(unprocess=False)

```

png


```

CPU times: user 1.58 s, sys: 300 ms, total: 1.88 s
Wall time: 2.11 s

```

## Training

We will be retraining a Google InceptionV3 model, pretrained on ImageNet. The neural network architecture is shown below.





```

1. %%time
2. from keras.models import Sequential, Model
3. from keras.layers import Dense, Dropout, Activation, Flatten
4. from keras.layers import Convolution2D, MaxPooling2D, ZeroPadding2D, GlobalAveragePooling2D
5. from keras.layers.normalization import BatchNormalization
6. from keras.preprocessing.image import ImageDataGenerator
7. from keras.callbacks import ModelCheckpoint, CSVLogger, LearningRateScheduler
8. from keras.optimizers import SGD
9. from keras.regularizers import l2
10. import keras.backend as K
11. import math
12.
13. K.clear_session()
14.
15. base_model = InceptionV3(weights='imagenet', include_top=False, input_tensor=X_train)
16. x = base_model.output
17. x = AveragePooling2D(pool_size=(8, 8))(x)
18. x = Dropout(.4)(x)
19. x = Flatten()(x)
20. predictions = Dense(n_classes, init='glorot_uniform', W_regularizer=l2(.0005))(x)
21.
22. model = Model(input=base_model.input, output=predictions)
23.
24. opt = SGD(lr=.01, momentum=.9)
25. model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
26.
27. checkpointer = ModelCheckpoint(filepath='model14.{epoch:02d}-{val_loss:.2f}.h5', save_best_only=True)
28. csv_logger = CSVLogger('model14.log')
29.
30. def schedule(epoch):
31.     if epoch < 15:
32.         return .01
33.     elif epoch < 28:
34.         return .002
35.     else:
36.         return .0004
37. lr_scheduler = LearningRateScheduler(schedule)
38.
39. model.fit_generator(train_generator,
40.                    validation_data=test_generator,
41.                    nb_val_samples=X_test.shape[0],
42.                    samples_per_epoch=X_train.shape[0],
43.                    nb_epoch=32,
44.                    verbose=2,
45.                    callbacks=[lr_scheduler, csv_logger, checkpointer])

```

Epoch 1/32

Epoch 00000: val\_loss improved from inf to 3.37355, saving model to model4.00-3.37355  
1342s - loss: 4.2541 - acc: 0.0810 - val\_loss: 3.3736 - val\_acc: 0.2010

Epoch 2/32

Epoch 00001: val\_loss improved from 3.37355 to 2.36625, saving model to model4.01-2.36625  
1329s - loss: 2.9745 - acc: 0.3075 - val\_loss: 2.3662 - val\_acc: 0.4071

Epoch 3/32

Epoch 00002: val\_loss improved from 2.36625 to 1.79355, saving model to model4.02-1.79355  
1329s - loss: 2.3080 - acc: 0.4539 - val\_loss: 1.7935 - val\_acc: 0.5338

Epoch 4/32

Epoch 00003: val\_loss improved from 1.79355 to 1.48898, saving model to model4.03-1.48898  
1356s - loss: 2.0102 - acc: 0.5216 - val\_loss: 1.4890 - val\_acc: 0.6068

Epoch 5/32

Epoch 00004: val\_loss improved from 1.48898 to 1.34121, saving model to model4.04-1.34121  
1330s - loss: 1.8436 - acc: 0.5577 - val\_loss: 1.3412 - val\_acc: 0.6431

Epoch 6/32

Epoch 00005: val\_loss improved from 1.34121 to 1.22485, saving model to model4.05-1.22485  
1329s - loss: 1.7057 - acc: 0.5909 - val\_loss: 1.2248 - val\_acc: 0.6740

Epoch 7/32

Epoch 00006: val\_loss did not improve

1328s - loss: 1.5996 - acc: 0.6126 - val\_loss: 1.2310 - val\_acc: 0.6716

Epoch 8/32

Epoch 00007: val\_loss improved from 1.22485 to 1.11248, saving model to model4.07-1.11248  
1331s - loss: 1.5148 - acc: 0.6314 - val\_loss: 1.1125 - val\_acc: 0.7022

Epoch 9/32

Epoch 00008: val\_loss improved from 1.11248 to 1.07145, saving model to model4.08-1.07145  
1331s - loss: 1.4395 - acc: 0.6506 - val\_loss: 1.0714 - val\_acc: 0.7095

Epoch 10/32

Epoch 00009: val\_loss improved from 1.07145 to 1.05129, saving model to model4.09-1.05129  
1333s - loss: 1.3900 - acc: 0.6637 - val\_loss: 1.0513 - val\_acc: 0.7181

Epoch 11/32

Epoch 00010: val\_loss improved from 1.05129 to 1.03356, saving model to model4.10-1.03356  
1331s - loss: 1.3316 - acc: 0.6780 - val\_loss: 1.0336 - val\_acc: 0.7250

Epoch 12/32

Epoch 00011: val\_loss improved from 1.03356 to 1.00622, saving model to model4.11-1.00622  
1331s - loss: 1.2850 - acc: 0.6893 - val\_loss: 1.0062 - val\_acc: 0.7275

Epoch 13/32

Epoch 00012: val\_loss improved from 1.00622 to 0.94016, saving model to model4.12-0.94016  
1330s - loss: 1.2325 - acc: 0.7003 - val\_loss: 0.9402 - val\_acc: 0.7461

Epoch 14/32

Epoch 00013: val\_loss did not improve

1330s - loss: 1.1970 - acc: 0.7086 - val\_loss: 0.9461 - val\_acc: 0.7453

Epoch 15/32

Epoch 00014: val\_loss did not improve

1329s - loss: 1.1683 - acc: 0.7154 - val\_loss: 0.9691 - val\_acc: 0.7396

Epoch 16/32

Epoch 00015: val\_loss improved from 0.94016 to 0.71776, saving model to model4.15-0.71776  
1329s - loss: 0.9398 - acc: 0.7724 - val\_loss: 0.7178 - val\_acc: 0.8055

Epoch 17/32

Epoch 00016: val\_loss improved from 0.71776 to 0.70245, saving model to model4.16-0.70245  
1329s - loss: 0.8591 - acc: 0.7916 - val\_loss: 0.7025 - val\_acc: 0.8069


```

Epoch 18/32
Epoch 00017: val_loss did not improve
1327s - loss: 0.8238 - acc: 0.8023 - val_loss: 0.7093 - val_acc: 0.8053
Epoch 19/32
Epoch 00018: val_loss did not improve
1327s - loss: 0.7947 - acc: 0.8093 - val_loss: 0.7048 - val_acc: 0.8059
Epoch 20/32
Epoch 00019: val_loss did not improve
1327s - loss: 0.7713 - acc: 0.8143 - val_loss: 0.7097 - val_acc: 0.8061
Epoch 21/32
Epoch 00020: val_loss improved from 0.70245 to 0.69545, saving model to model4.20-0
1329s - loss: 0.7458 - acc: 0.8195 - val_loss: 0.6955 - val_acc: 0.8104
Epoch 22/32
Epoch 00021: val_loss did not improve
1328s - loss: 0.7282 - acc: 0.8232 - val_loss: 0.6977 - val_acc: 0.8119
Epoch 23/32
Epoch 00022: val_loss improved from 0.69545 to 0.69190, saving model to model4.22-0
1328s - loss: 0.7114 - acc: 0.8284 - val_loss: 0.6919 - val_acc: 0.8150
Epoch 24/32
Epoch 00023: val_loss did not improve
1325s - loss: 0.6983 - acc: 0.8311 - val_loss: 0.7002 - val_acc: 0.8116
Epoch 25/32
Epoch 00024: val_loss did not improve
1330s - loss: 0.6719 - acc: 0.8381 - val_loss: 0.7031 - val_acc: 0.8112
Epoch 26/32
Epoch 00025: val_loss did not improve
1382s - loss: 0.6607 - acc: 0.8407 - val_loss: 0.7115 - val_acc: 0.8083
Epoch 27/32
Epoch 00026: val_loss did not improve
1330s - loss: 0.6479 - acc: 0.8439 - val_loss: 0.7037 - val_acc: 0.8126
Epoch 28/32
Epoch 00027: val_loss did not improve
1328s - loss: 0.6292 - acc: 0.8478 - val_loss: 0.7122 - val_acc: 0.8086
Epoch 29/32
Epoch 00028: val_loss improved from 0.69190 to 0.68908, saving model to model4.28-0
1330s - loss: 0.5983 - acc: 0.8580 - val_loss: 0.6891 - val_acc: 0.8165
Epoch 30/32
Epoch 00029: val_loss improved from 0.68908 to 0.68740, saving model to model4.29-0
1330s - loss: 0.5817 - acc: 0.8612 - val_loss: 0.6874 - val_acc: 0.8149
Epoch 31/32
Epoch 00030: val_loss did not improve
1328s - loss: 0.5729 - acc: 0.8642 - val_loss: 0.6912 - val_acc: 0.8143
Epoch 32/32
Epoch 00031: val_loss did not improve
1329s - loss: 0.5638 - acc: 0.8663 - val_loss: 0.6895 - val_acc: 0.8159
CPU times: user 8h 49min 20s, sys: 1h 55min 54s, total: 10h 45min 14s
Wall time: 11h 51min 18s

```

At this point, we are seeing up to 81.65 single crop Top-1 accuracy on the test set. We can continue to train the model at an even slower learning rate to see if it improves more.

My initial experiments used more modern optimizers such as Adam and AdaDelta, along with higher learning rates. I was stuck for a while below 80% accuracy before I decided to follow the literature more closely and use Stochastic Gradient Descent (SGD) with a quickly decreasing learning schedule. When we are searching through the multidimensional surface, sometimes going slower goes a long way.



Due to some instability with my multiprocessing code, sometimes I need to restart the notebook, load the latest model, then continue training.

```

1. %%time
2. from keras.models import Sequential, Model, load_model
3. from keras.layers import Dense, Dropout, Activation, Flatten
4. from keras.layers import Convolution2D, MaxPooling2D, ZeroPadding2D, GlobalAveragePooling2D
5. from keras.layers.normalization import BatchNormalization
6. from keras.preprocessing.image import ImageDataGenerator
7. from keras.callbacks import ModelCheckpoint, CSVLogger, LearningRateScheduler
8. from keras.optimizers import SGD
9. from keras.regularizers import l2
10. import keras.backend as K
11. import math
12.
13. model = load_model(filepath='./model4.29-0.69.hdf5')
14.
15. opt = SGD(lr=.01, momentum=.9)
16. model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
17.
18. checkpointer = ModelCheckpoint(filepath='model4b.{epoch:02d}-{val_loss:.2f}.hdf5',
19. csv_logger = CSVLogger('model4b.log'))
20.
21. def schedule(epoch):
22.     if epoch < 10:
23.         return .00008
24.     elif epoch < 20:
25.         return .00016
26.     else:
27.         return .000032
28.
29. lr_scheduler = LearningRateScheduler(schedule)
30.
31. model.fit_generator(train_generator,
32.                     validation_data=test_generator,
33.                     nb_val_samples=X_test.shape[0],
34.                     samples_per_epoch=X_train.shape[0],
35.                     nb_epoch=32,
36.                     verbose=2,
37.                     callbacks=[lr_scheduler, csv_logger, checkpointer])

```

## Model Evaluation

At this point, we should have multiple trained models saved to disk. We can go through them and use the `load_model` function to load the model with the lowest loss / highest accuracy.

```

1. %%time
2. #model = load_model(filepath='./model4.29-0.69.hdf5') # 86.8039 10-crop Top-1
3. model = load_model(filepath='./model4b.10-0.68.hdf5') # 86.9703

```

```

CPU times: user 36.4 s, sys: 1.11 s, total: 37.5 s
Wall time: 36.5 s

```

We also want to evaluate the test set using multiple crops. This can yield an accuracy boost of 5% compared to single crop evaluation. It is common to use the following crops: Upper Left, Upper Right, Lower Left, Lower Right, Center. We also take the same crops on the image flipped left to right, creating a total of 10 crops.

In addition, we want to return the top-N predictions for each crop in order to calculate Top-5 accuracy, for instance.

```
1. def center_crop(x, center_crop_size, **kwargs):
2.     centerw, centerh = x.shape[0]//2, x.shape[1]//2
3.     halfw, halfh = center_crop_size[0]//2, center_crop_size[1]//2
4.     return x[centerw-halfw:centerw+halfw+1, centerh-halfh:centerh+halfh+1, :]
```

```

1. def predict_10_crop(img, ix, top_n=5, plot=False, preprocess=True, debug=False):
2.     flipped_X = np.fliplr(img)
3.     crops = [
4.         img[:299,:299, :], # Upper Left
5.         img[:299, img.shape[1]-299:, :], # Upper Right
6.         img[img.shape[0]-299:, :299, :], # Lower Left
7.         img[img.shape[0]-299:, img.shape[1]-299:, :], # Lower Right
8.         center_crop(img, (299, 299)),
9.
10.        flipped_X[:299,:299, :],
11.        flipped_X[:299, flipped_X.shape[1]-299:, :],
12.        flipped_X[flipped_X.shape[0]-299:, :299, :],
13.        flipped_X[flipped_X.shape[0]-299:, flipped_X.shape[1]-299:, :],
14.        center_crop(flipped_X, (299, 299))
15.    ]
16.    if preprocess:
17.        crops = [preprocess_input(x.astype('float32')) for x in crops]
18.
19.    if plot:
20.        fig, ax = plt.subplots(2, 5, figsize=(10, 4))
21.        ax[0][0].imshow(crops[0])
22.        ax[0][1].imshow(crops[1])
23.        ax[0][2].imshow(crops[2])
24.        ax[0][3].imshow(crops[3])
25.        ax[0][4].imshow(crops[4])
26.        ax[1][0].imshow(crops[5])
27.        ax[1][1].imshow(crops[6])
28.        ax[1][2].imshow(crops[7])
29.        ax[1][3].imshow(crops[8])
30.        ax[1][4].imshow(crops[9])
31.
32.    y_pred = model.predict(np.array(crops))
33.    preds = np.argmax(y_pred, axis=1)
34.    top_n_preds= np.argpartition(y_pred, -top_n)[:,-top_n:]
35.    if debug:
36.        print('Top-1 Predicted:', preds)
37.        print('Top-5 Predicted:', top_n_preds)
38.        print('True Label:', y_test[ix])
39.    return preds, top_n_preds
40.
41.
42. ix = 13001
43. predict_10_crop(X_test[ix], ix, top_n=5, plot=True, preprocess=False, debug=True)

```

```
Top-1 Predicted: [74 74 74 74 74 74 74 74 74 74]
```

```
Top-5 Predicted: [[33 97 37 39 74]
```

```
[28 52 37 39 74]
```

```
[73 39 52 37 74]
```

```
[35 33 37 39 74]
```

```
[35 33 37 39 74]
```

```
[35 33 37 39 74]
```

```
[35 33 37 39 74]
```

```
[97 37 73 39 74]
```

```
[73 52 37 39 74]
```

```
[34 35 33 39 74]]
```

```
True Label: 88
```

```
(array([74, 74, 74, 74, 74, 74, 74, 74, 74, 74]), array([[33, 97, 37, 39, 74],  
                [28, 52, 37, 39, 74],  
                [73, 39, 52, 37, 74],  
                [35, 33, 37, 39, 74],  
                [35, 33, 37, 39, 74],  
                [35, 33, 37, 39, 74],  
                [35, 33, 37, 39, 74],  
                [97, 37, 73, 39, 74],  
                [73, 52, 37, 39, 74],  
                [34, 35, 33, 39, 74]]))
```

png

We also need to preprocess the images for the Inception model:

1. `ix = 13001`
2. `predict_10_crop(X_test[ix], ix, top_n=5, plot=True, preprocess=True, debug=Tr`



```
Top-1 Predicted: [51 51 88 88 88 51 51 88 88 88]
```

```
Top-5 Predicted: [[18 79 51 13 48]
```

```
[48 79 11 55 51]
```

```
[79 93 81 37 88]
```

```
[51 86 93 81 88]
```

```
[11 79 51 81 88]
```

```
[19 79 51 56 13]
```

```
[11 88 48 51 13]
```

```
[37 93 86 88 81]
```

```
[37 79 93 88 81]
```

```
[84 81 11 79 88]]
```

```
True Label: 88
```

```
(array([51, 51, 88, 88, 88, 51, 51, 88, 88, 88]), array([[18, 79, 51, 13, 48],  
               [48, 79, 11, 55, 51],  
               [79, 93, 81, 37, 88],  
               [51, 86, 93, 81, 88],  
               [11, 79, 51, 81, 88],  
               [19, 79, 51, 56, 13],  
               [11, 88, 48, 51, 13],  
               [37, 93, 86, 88, 81],  
               [37, 79, 93, 88, 81],  
               [84, 81, 11, 79, 88]]))
```

png

Now we create crops for each item in the test set and get the predictions. This is a slow process at the moment as I am not taking advantage of multiprocessing or other types of parallelism.

```
1. %%time  
2. preds_10_crop = {}  
3. for ix in range(len(X_test)):  
4.     if ix % 1000 == 0:  
5.         print(ix)  
6.         preds_10_crop[ix] = predict_10_crop(X_test[ix], ix)
```

```
0
1000
2000
3000
4000
5000
6000
7000
8000
9000
10000
11000
12000
13000
14000
15000
16000
17000
18000
19000
20000
21000
22000
23000
24000
25000
CPU times: user 50min 3s, sys: 5min 13s, total: 55min 16s
Wall time: 31min 28s
```

We now have a set of 10 predictions for each image. Using a histogram, I'm able to see how the # of unique predictions for each image are distributed.

```
1. preds_uniq = {k: np.unique(v[0]) for k, v in preds_10_crop.items()}
2. preds_hist = np.array([len(x) for x in preds_uniq.values()])
3.
4. plt.hist(preds_hist, bins=11)
5. plt.title('Number of unique predictions per image')
```

```
<matplotlib.text.Text at 0x7fe30c3daa20>
```

png

Let's create a dictionary to map test item index to its top-1 / top-5 predictions.

```
1. preds_top_1 = {k: collections.Counter(v[0]).most_common(1) for k, v in preds_
2.
3. top_5_per_ix = {k: collections.Counter(preds_10_crop[k][1].reshape(-1)).most_
4.                 for k, v in preds_10_crop.items()}
5. preds_top_5 = {k: [y[0] for y in v] for k, v in top_5_per_ix.items()}
```

```
1. %%time
2. right_counter = 0
3. for i in range(len(y_test)):
4.     guess, actual = preds_top_1[i][0][0], y_test[i]
5.     if guess == actual:
6.         right_counter += 1
7.
8. print('Top-1 Accuracy, 10-Crop: {0:.2f}%'.format(right_counter / len(y_test))
```

Top-1 Accuracy, 10-Crop: 86.97%  
CPU times: user 28 ms, sys: 0 ns, total: 28 ms  
Wall time: 27.3 ms

```
1. %%time
2. top_5_counter = 0
3. for i in range(len(y_test)):
4.     guesses, actual = preds_top_5[i], y_test[i]
5.     if actual in guesses:
6.         top_5_counter += 1
7.
8. print('Top-5 Accuracy, 10-Crop: {0:.2f}%'.format(top_5_counter / len(y_test))
```

Top-5 Accuracy, 10-Crop: 97.42%  
CPU times: user 28 ms, sys: 0 ns, total: 28 ms  
Wall time: 27 ms

## Results Visualization

```
1. y_pred = [x[0][0] for x in preds_top_1.values()]
```

```

1. @interact(page=[0, int(len(X_test)/20)])
2. def show_images_prediction(page=0):
3.     page_size = 20
4.     nrows = 4
5.     ncols = 5
6.     fig, axes = plt.subplots(nrows=nrows, ncols=ncols, figsize=(12, 12))
7.     fig.set_size_inches(12, 8)
8.     #fig.tight_layout()
9.     #imgs = np.random.choice((y_all == n_class).nonzero()[0], nrows * ncols)
10.    start_i = page * page_size
11.    for i, ax in enumerate(axes.flat):
12.        im = ax.imshow(X_test[i+start_i])
13.        ax.set_axis_off()
14.        ax.title.set_visible(False)
15.        ax.xaxis.set_ticks([])
16.        ax.yaxis.set_ticks([])
17.        for spine in ax.spines.values():
18.            spine.set_visible(False)
19.        predicted = ix_to_class[y_pred[i+start_i]]
20.        match = predicted == ix_to_class[y_test[start_i + i]]
21.        ec = (1, .5, .5)
22.        fc = (1, .8, .8)
23.        if match:
24.            ec = (0, .6, .1)
25.            fc = (0, .7, .2)
26.            # predicted label
27.            ax.text(0, 400, 'P: ' + predicted, size=10, rotation=0,
28.                ha="left", va="top",
29.                bbox=dict(boxstyle="round",
30.                    ec=ec,
31.                    fc=fc,
32.                )
33.            )
34.        if not match:
35.            # true label
36.            ax.text(0, 480, 'A: ' + ix_to_class[y_test[start_i + i]], size=10,
37.                ha="left", va="top",
38.                bbox=dict(boxstyle="round",
39.                    ec=ec,
40.                    fc=fc,
41.                )
42.            )
43.    plt.subplots_adjust(left=0, wspace=1, hspace=0)
44.    plt.show()

```

png

A confusion matrix will plot each class label and how many times it was correctly labeled vs. the other times it was incorrectly labeled as a different class.



Confusion matrix, without normalization

```
[[179  0  4 ...,  2  0  5]
 [  0 218  0 ...,  0  0  0]
 [  4  0 228 ...,  1  0  0]
 ...,
 [  0  0  0 ..., 212  0  1]
 [  0  0  0 ...,  0 208  0]
 [  0  0  0 ...,  0  0 224]]
```

png

CPU times: user 16.4 s, sys: 1.22 s, total: 17.6 s  
Wall time: 16.4 s

We want to see if the accuracy was consistent across all classes, or if some classes were much easier / harder to label than others. According to our plot, a few classes were outliers in terms of being much more difficult to label correctly.

```
1. corrects = collections.defaultdict(int)
2. incorrects = collections.defaultdict(int)
3. for (pred, actual) in zip(y_pred, y_test):
4.     if pred == actual:
5.         corrects[actual] += 1
6.     else:
7.         incorrects[actual] += 1
8.
9. class_accuracies = {}
10. for ix in range(101):
11.     class_accuracies[ix] = corrects[ix]/250
12.
13. plt.hist(list(class_accuracies.values()), bins=20)
14. plt.title('Accuracy by Class histogram')
```

<matplotlib.text.Text at 0x7fe2d5d4f860>

png

```
1. sorted_class_accuracies = sorted(class_accuracies.items(), key=lambda x: -x[1])
2. [(ix_to_class[c[0]], c[1]) for c in sorted_class_accuracies]
```

```
[('edamame', 0.996),
 ('hot_and_sour_soup', 0.964),
 ('oysters', 0.964),
 ('seaweed_salad', 0.96),
 ('macarons', 0.956),
 ('pad_thai', 0.956),
 ('spaghetti_bolognese', 0.956),
 ('french_fries', 0.952),
 ('frozen_yogurt', 0.952),
 ('takoyaki', 0.952),
 ('spaghetti_carbonara', 0.948),
 ('clam_chowder', 0.944),
 ('deviled_eggs', 0.944),
 ('churros', 0.94),
 ('miso_soup', 0.94),
 ('creme_brulee', 0.936),
 ('pho', 0.936),
 ('cannoli', 0.932),
 ('guacamole', 0.932),
 ('mussels', 0.932),
 ('sashimi', 0.932),
 ('caesar_salad', 0.928),
 ('lobster_roll_sandwich', 0.928),
 ('bibimbap', 0.924),
 ('cup_cakes', 0.924),
 ('dumplings', 0.924),
 ('ramen', 0.924),
 ('beef_carpaccio', 0.92),
 ('eggs_benedict', 0.92),
 ('pancakes', 0.92),
 ('red_velvet_cake', 0.92),
 ('beignets', 0.916),
 ('club_sandwich', 0.916),
 ('escargots', 0.916),
 ('french_onion_soup', 0.916),
 ('onion_rings', 0.916),
 ('baklava', 0.912),
 ('croque_madame', 0.912),
 ('fish_and_chips', 0.908),
 ('poutine', 0.908),
 ('cheese_plate', 0.904),
 ('chicken_wings', 0.904),
 ('fried_rice', 0.904),
 ('sushi', 0.904),
 ('fried_calamari', 0.9),
 ('pulled_pork_sandwich', 0.896),
 ('waffles', 0.896),
 ('crab_cakes', 0.892),
 ('gyoza', 0.892),
 ('paella', 0.892),
 ('caprese_salad', 0.888),
```

```
('lobster_bisque', 0.888),
('peking_duck', 0.888),
('pizza', 0.888),
('greek_salad', 0.88),
('hot_dog', 0.88),
('samosa', 0.88),
('donuts', 0.876),
('spring_rolls', 0.876),
('baby_back_ribs', 0.872),
('strawberry_shortcake', 0.872),
('shrimp_and_grits', 0.868),
('tacos', 0.86),
('beef_tartare', 0.856),
('prime_rib', 0.856),
('chicken_quesadilla', 0.852),
('hummus', 0.852),
('grilled_salmon', 0.848),
('tiramisu', 0.848),
('macaroni_and_cheese', 0.844),
('carrot_cake', 0.836),
('nachos', 0.836),
('falafel', 0.832),
('tuna_tartare', 0.832),
('panna_cotta', 0.828),
('bruschetta', 0.824),
('grilled_cheese_sandwich', 0.824),
('risotto', 0.812),
('french_toast', 0.808),
('gnocchi', 0.808),
('garlic_bread', 0.804),
('breakfast_burrito', 0.8),
('beet_salad', 0.796),
('hamburger', 0.796),
('cheesecake', 0.792),
('lasagna', 0.792),
('ceviche', 0.784),
('chicken_curry', 0.784),
('omelette', 0.784),
('scallops', 0.784),
('chocolate_cake', 0.78),
('huevos_rancheros', 0.78),
('ravioli', 0.776),
('ice_cream', 0.764),
('bread_pudding', 0.748),
('foie_gras', 0.72),
('apple_pie', 0.716),
('filet_mignon', 0.716),
('chocolate_mousse', 0.7),
('pork_chop', 0.676),
('steak', 0.576)]
```



## Interactive Classification

## Predicting from a local file

```
1. pic_path = '/home/stratospark/Downloads/soup.jpg'
2. pic = img.imread(pic_path)
3. preds = predict_10_crop(np.array(pic), 0)[0]
4. best_pred = collections.Counter(preds).most_common(1)[0][0]
5. print(ix_to_class[best_pred])
6. plt.imshow(pic)
```

french\_onion\_soup

```
<matplotlib.image.AxesImage at 0x7fe2d59eb5c0>
```

png

## Predicting from an image on the Internet

```
1. import urllib.request
2.
3. @interact
4. def predict_remote_image(url='http://themodelhouse.tv/wp-content/uploads/2016
5.     with urllib.request.urlopen(url) as f:
6.         pic = plt.imread(f, format='jpg')
7.         preds = predict_10_crop(np.array(pic), 0)[0]
8.         best_pred = collections.Counter(preds).most_common(1)[0][0]
9.         print(ix_to_class[best_pred])
10.         plt.imshow(pic)
```

hummus

png

## Keras.js Export

```
1. with open('model.json', 'w') as f:
2.     f.write(model.to_json())
```

```
1. import json
2.
3. json.dumps(ix_to_class)
```

```
'{"0": "apple_pie", "1": "baby_back_ribs", "2": "baklava", "3": "beef_carpaccio", '}
```